AI Impact Summit 2026

Technical Conference Report
Date of Attendance: 17 February 2026

1. From Policy to Practice: Governing AI for Global Impact

Panel Participants

- ***NASSCOM***
- ***Gail Kent – Google***
- ***Ivana – Wipro***
- ***Jules – Future of Privacy Forum (FPF)***

Session Overview

The session focused on translating AI governance frameworks into practical implementation strategies. The discussion emphasized the evolution of AI systems toward agentic architectures capable of autonomous reasoning, tool usage, and multi-step task execution.

Key Technical Insights

- Governance must be embedded into AI system design rather than applied post-deployment.
- Agentic systems require structured oversight, transparency, explainability, and regulatory alignment.
- Google aims to democratize agent development through drag-and-drop, low-code platforms.
- Frameworks such as Microsoft Studio and Google Agent Development Kit are accelerating enterprise adoption.

Ivana (Wipro) highlighted the strategic role of consulting organizations in AI transformation. While product companies build foundational platforms, consulting firms translate business requirements into deployable AI architectures. The value lies in selecting appropriate approaches (Agentic systems, RAG pipelines, hybrid AI models), ensuring integration with legacy systems, and delivering measurable ROI.

Future Outlook

The panel discussed expansion areas for Agentic AI and Retrieval-Augmented Generation (RAG) systems, particularly in enterprise automation, domain-specific knowledge systems, and scalable governance frameworks.

2. Powering AI with Quantum

Panel Participants

- ***Amit Singhee – IBM***
- ***Gopal Ranganathan – Chicago-based Startup***

Others

Technical Discussion

The session explored the convergence of Artificial Intelligence and Quantum Computing. Currently, AI workloads run on High-Performance Computing (HPC) systems utilizing CPUs and GPUs. The discussion introduced Quantum Processing Units (QPUs) as the next evolution in computational capability.

Amit Singhee explained the distinction between Classical Neural Networks and Quantum Neural Networks (QNNs). While classical models rely on deterministic weight optimization, QNNs leverage quantum superposition and entanglement to potentially solve complex optimization problems more efficiently.

IBM is actively developing QPU-based AI frameworks and enabling the transformation of PyTorch and TensorFlow models for quantum environments. IBM is also contributing to the Quantum Valley initiative in Amaravati.

During a one-to-one discussion, Amit recommended Qiskit as a starting platform for Quantum + AI development. IBM provides access to a 16-qubit quantum system for experimentation and execution of quantum programs.

Additional topics covered by Gopal Ranganathan included Quantum Circuits, Quantum Optimization, Quantum Algorithms, Quantum Sensing, and Quantum Processing architectures.

3. India AI Stack

Panel Participants

• ***Col. Amit Mehta***

Others

Strategic Focus Areas

• Application Layer

• Hardware Infrastructure

• Implementation Strategy

The discussion focused on national defense requirements for AI deployment. **Col. Amit** emphasized the need for high-accuracy, locally deployed AI models to meet operational security standards.

The Indian Army is developing and fine-tuning Large Language Models (LLMs) on defense-specific datasets for localized usage. A hybrid Human + AI decision-making framework is being explored to enhance situational awareness.

Additional use cases include satellite image processing, intelligence analysis, and secure on-premise AI systems.

4. Putting AI to Work: Solving Productivity Challenges through Faster Adoption

Panel Participants

- ***David Zapolsky – Amazon***
- ***Kanishka Narayan – Ministry of AI, UK***
- ***Kevin Allison – Minerva Technology Futures***
- ***Shubhi Agrawal – LocoBuzz***

Key Highlights

The session emphasized accelerating AI adoption to drive measurable productivity gains across industries.

LocoBuzz, an Indian startup, has partnered with Amazon Bedrock to build customized agent-based solutions for enterprise clients. Amazon Bedrock enables scalable model deployment and managed AI infrastructure.

David Zapolsky explained the usability and integration features of Bedrock, highlighting its flexibility for enterprise-grade agent development.

Kanishka Narayan discussed public-sector AI implementation strategies, focusing on responsible AI deployment for citizen services and governance.

The session concluded with a consensus that faster AI adoption requires strong governance, scalable infrastructure, industry collaboration, and practical use-case driven implementation.

Exhibition Exploration Report

Exhibition Exploration Report
Date: 18 February 2026

1. Arduino & Qualcomm Edge AI Development

During the exhibition, I visited the Arduino booth showcasing next-generation boards developed in partnership with Qualcomm. These boards are designed to enable edge computing capabilities with integrated AI/ML acceleration.

A live demonstration highlighted real-time Face Detection running directly on the Arduino board, demonstrating on-device inference without reliance on cloud infrastructure. The platform supports seamless integration with cameras and multiple sensors, making prototyping efficient for embedded AI applications.

In discussion with the technical team, they explained the architectural differences between this board and Raspberry Pi. Unlike Raspberry Pi, which is a general-purpose single-board computer, the new Arduino-Qualcomm board focuses on low-power edge AI inference, optimized ML acceleration, and embedded deployment scenarios.

Arduino Uno Q

2. Intel Workshop – Agentic Stock Market & Trading Systems

I attended an Intel-led workshop focused on building agentic systems for stock market analysis and automated trading.

The solution architecture leveraged the Qwen 2.5 large language model integrated with the Zerodha API for real-time trading data access. Intel utilized the Intel AI Super Builder framework to develop and orchestrate autonomous trading agents.

The workshop demonstrated how LLM-based agents can analyze market signals, interpret structured financial data, and execute rule-based trading strategies. The session highlighted the importance of latency optimization, API orchestration, and risk governance in financial AI systems.

3. Meta – Edge NLP & Real-Time Interaction

At the Meta exhibition area, I explored implementations of Natural Language Processing (NLP) models deployed on edge devices for real-time interaction.

The demonstration showed how optimized NLP models can run locally on compact hardware to enable low-latency conversational experiences. The architecture reduces cloud dependency while ensuring faster inference and enhanced privacy.

4. Tyrone Systems – High-Performance AI Servers

I visited the Tyrone Systems booth and explored their enterprise-grade AI server infrastructure, including the B200 and B300 server configurations.

These servers are engineered for large-scale AI model training, testing, and inference workloads. They support computationally intensive models such as Convolutional Neural Networks (CNN), Recurrent Neural Networks (RNN), Artificial Neural Networks (ANN), and Large Language Models (LLMs).

The infrastructure design focuses on GPU acceleration, high-throughput data pipelines, scalable memory architecture, and optimized cooling systems for sustained performance.

5. QNU Labs – Quantum Security Platform

I explored QNU Labs Quantum Security Platform, which delivers quantum-safe encryption solutions via a Software-as-a-Service (SaaS) model.

Their platform leverages Quantum Key Distribution (QKD) principles, implementing secure cryptographic key exchange mechanisms often illustrated using the 'Alice and Bob' protocol model.

The solution is designed to protect enterprise software systems from external threats by ensuring quantum-resistant encryption standards. The platform integrates with existing network infrastructure to enhance cybersecurity resilience against future quantum attacks.